

# Generic Application-Server Crates

## Built With Async Rust

Brendan Nolan

August 22, 2023

- ▶ Provide an *application-server* crate and a *client* crate, allowing users to request jobs which are too heavy for their local machines.
- ▶ Example use case: scientists in a wetlab, who must regularly run heavy models from their bench laptops.

- ▶ Provide an *application-server* crate and a *client* crate, allowing users to request jobs which are too heavy for their local machines.
- ▶ Example use case: scientists in a wetlab, who must regularly run heavy models from their bench laptops.

## How A User Creates A Custom Server

A user who wants a server that provides weather status must only do the following:

```
let (tx_shutdown, rx_shutdown) = watch::channel(());
let server_task = tokio::spawn(server_runner::run_server(
    "<IP Address>:<Port>",
    weather_provider, // User's type — must implement RequestProcessor
    ShutdownListener::new(rx_shutdown),
));
```

## Example Application

Weather Type

Wind | Rain | Sun

Counties

Wexford

Cork

Forecast

Rain

Rain forecast for Wexford is 0.9267098

## Must-Have Features

For crate users

- ▶ Easily customise server data processing, e.g.:
  - ▶ train ML algorithms
  - ▶ run planetary models
  - ▶ run structural analysis of timberframe houses
- ▶ Allow client to choose between
  - ▶ call-and-response model
  - ▶ sending and receiving in any order
- ▶ Remain responsive as number of clients connecting to a single server grows

## Must-Have Features

### For crate users

- ▶ Easily customise server data processing, e.g.:
  - ▶ train ML algorithms
  - ▶ run planetary models
  - ▶ run structural analysis of timberframe houses
- ▶ Allow client to choose between
  - ▶ call-and-response model
  - ▶ sending and receiving in any order
- ▶ Remain responsive as number of clients connecting to a single server grows

## Must-Have Features

### For crate users

- ▶ Easily customise server data processing, e.g.:
  - ▶ train ML algorithms
  - ▶ run planetary models
  - ▶ run structural analysis of timberframe houses
- ▶ Allow client to choose between
  - ▶ call-and-response model
  - ▶ sending and receiving in any order
- ▶ Remain responsive as number of clients connecting to a single server grows



For crate developers

- ▶ Do not waste resources by blocking threads while waiting for I/O
- ▶ (More insidious than the above) do not make asynchronous work wait for other asynchronous work unnecessarily

### For crate developers

- ▶ Do not waste resources by blocking threads while waiting for I/O
- ▶ (More insidious than the above) do not make asynchronous work wait for other asynchronous work unnecessarily

### Criteria for chosen technology

- ▶ Has well documented library(s) for async I/O
- ▶ Eliminates as many bugs as possible as early as possible (strong bias towards compiled languages)

### Criteria for chosen technology

- ▶ Has well documented library(s) for async I/O
- ▶ Eliminates as many bugs as possible as early as possible (strong bias towards compiled languages)

### *C++: boost.asio*

- ▶ Pros

- ▶ *boost.asio* is a mature library

- ▶ Cons

- ▶ awkward call-back mechanism
  - ▶ many object-lifetime footguns (use-after-free bugs)
  - ▶ static polymorphism extremely awkward to use

### *C++: boost.asio*

- ▶ Pros

- ▶ *boost.asio* is a mature library

- ▶ Cons

- ▶ awkward call-back mechanism
  - ▶ many object-lifetime footguns (use-after-free bugs)
  - ▶ static polymorphism extremely awkward to use

### *Rust: tokio*

- ▶ Cons

- ▶ async I/O ecosystem in *Rust* is still young

- ▶ Pros

- ▶ *async/await* syntax is far more ergonomic than passing callbacks around
  - ▶ typesystem eliminates use-after-free bugs
  - ▶ static polymorphism is user friendly

### *Rust: tokio*

- ▶ Cons

- ▶ *async I/O ecosystem in Rust is still young*

- ▶ Pros

- ▶ *async/await* syntax is far more ergonomic than passing callbacks around
  - ▶ typesystem eliminates use-after-free bugs
  - ▶ static polymorphism is user friendly



Winner: *Rust* with *tokio*

- Customization point should allow state (e.g. in order to cache results of previous model runs)

```
pub trait RequestProcessor<Req, Resp> {  
    fn process(&self, request: &Req) -> Resp;  
}
```

- ▶ Support serialization by *serde* only
- ▶ Make three crates (for client, server, and I/O utilities) rather than combining all into one crate

- ▶ Support serialization by *serde* only
- ▶ Make three crates (for client, server, and I/O utilities) rather than combining all into one crate

- ▶ Follow the Golang philosophy: "Do not communicate by sharing memory - share memory by communicating"
- ▶ Employ the *actor* pattern, recommended by Alice Ryhl (tokio maintainer)
- ▶ Make it as simple as possible to set up a custom server

- ▶ Follow the Golang philosophy: "Do not communicate by sharing memory - share memory by communicating"
- ▶ Employ the *actor* pattern, recommended by Alice Ryhl (tokio maintainer)
- ▶ Make it as simple as possible to set up a custom server

- ▶ Follow the Golang philosophy: "Do not communicate by sharing memory - share memory by communicating"
- ▶ Employ the *actor* pattern, recommended by Alice Ryhl (tokio maintainer)
- ▶ Make it as simple as possible to set up a custom server

## How The Server Handles Jobs

### My Implementation Of The Actor Pattern

(Some generic parameters and trait bounds removed for readability)

```
pub struct Command<Req, Resp> {
    pub data: Req,
    pub responder: oneshot::Sender<Resp>,
}

// The "Handle"
pub struct JobDispatcher<Req, Resp> {
    // Receiver owned by "Actor" (jobs task)
    tx: mpsc::Sender<Command<Req, Resp>>,
}

impl Clone for JobDispatcher<Req, Resp> { ... }

impl JobDispatcher<Req, Resp> {
    pub async fn dispatch_job(&self, data: Req) -> oneshot::Receiver<Resp> {
        // Create command from data
        // Send command to jobs task
        // Return receiver from which job result can be collected
    }
}
```



*server* crate:

- ▶ Implement unified and ergonomic way to handle shutdown signals
- ▶ Add organised logging (with support for various levels of verbosity)
- ▶ Consider compatability of serialized data across architectures, versions of this crate, versions of dependencies (look into protobuf). (May be difficult because input/output data is generic by design.)
- ▶ Think carefully about backpressure
- ▶ Ensure there are no partial-read vulnerabilities
- ▶ Investigate automatically figuring out channel bounds, rather than asking the user to supply them
- ▶ Automated testing

### *server* crate:

- ▶ Implement unified and ergonomic way to handle shutdown signals
- ▶ Add organised logging (with support for various levels of verbosity)
- ▶ Consider compatability of serialized data across architectures, versions of this crate, versions of dependencies (look into protobuf). (May be difficult because input/output data is generic by design.)
- ▶ Think carefully about backpressure
- ▶ Ensure there are no partial-read vulnerabilities
- ▶ Investigate automatically figuring out channel bounds, rather than asking the user to supply them
- ▶ Automated testing

### *server* crate:

- ▶ Implement unified and ergonomic way to handle shutdown signals
- ▶ Add organised logging (with support for various levels of verbosity)
- ▶ Consider compatability of serialized data across architectures, versions of this crate, versions of dependencies (look into protobuf). (May be difficult because input/output data is generic by design.)
- ▶ Think carefully about backpressure
- ▶ Ensure there are no partial-read vulnerabilities
- ▶ Investigate automatically figuring out channel bounds, rather than asking the user to supply them
- ▶ Automated testing

### *server crate:*

- ▶ Implement unified and ergonomic way to handle shutdown signals
- ▶ Add organised logging (with support for various levels of verbosity)
- ▶ Consider compatability of serialized data across architectures, versions of this crate, versions of dependencies (look into protobuf). (May be difficult because input/output data is generic by design.)
- ▶ **Think carefully about backpressure**
- ▶ Ensure there are no partial-read vulnerabilities
- ▶ Investigate automatically figuring out channel bounds, rather than asking the user to supply them
- ▶ Automated testing

### *server* crate:

- ▶ Implement unified and ergonomic way to handle shutdown signals
- ▶ Add organised logging (with support for various levels of verbosity)
- ▶ Consider compatability of serialized data across architectures, versions of this crate, versions of dependencies (look into protobuf). (May be difficult because input/output data is generic by design.)
- ▶ Think carefully about backpressure
- ▶ **Ensure there are no partial-read vulnerabilities**
- ▶ Investigate automatically figuring out channel bounds, rather than asking the user to supply them
- ▶ Automated testing

### *server* crate:

- ▶ Implement unified and ergonomic way to handle shutdown signals
- ▶ Add organised logging (with support for various levels of verbosity)
- ▶ Consider compatability of serialized data across architectures, versions of this crate, versions of dependencies (look into protobuf). (May be difficult because input/output data is generic by design.)
- ▶ Think carefully about backpressure
- ▶ Ensure there are no partial-read vulnerabilities
- ▶ Investigate automatically figuring out channel bounds, rather than asking the user to supply them
- ▶ Automated testing

### *server crate:*

- ▶ Implement unified and ergonomic way to handle shutdown signals
- ▶ Add organised logging (with support for various levels of verbosity)
- ▶ Consider compatability of serialized data across architectures, versions of this crate, versions of dependencies (look into protobuf). (May be difficult because input/output data is generic by design.)
- ▶ Think carefully about backpressure
- ▶ Ensure there are no partial-read vulnerabilities
- ▶ Investigate automatically figuring out channel bounds, rather than asking the user to supply them
- ▶ **Automated testing**

*client* crate:

- ▶ Not I/O bound - does it need async? Would it be enough to have a read thread and a write thread?